Segmentation

Lu Peng
School of Computer Science,
Beijing University of Posts and Telecommunications

Machine Vision Technology							
Semantic information				**Metric 3D information**			
Pixels	Segments	Images	Videos	Camera		Multi-view Geometry	
Convolutions Edges & Fitting Local features Texture	Segmentation Clustering	Recognition Detection	Motion Tracking	Camera Model	Camera Calibration	Epipolar Geometry	SFM
10	4	4	2	2	2	2	2

Image segmentation

Source： S. Lazebnik

The goals of segmentation

- Group together similar-looking pixels for efficiency of further processing
  - "Bottom-up" process
  - Unsupervised

"superpixels"

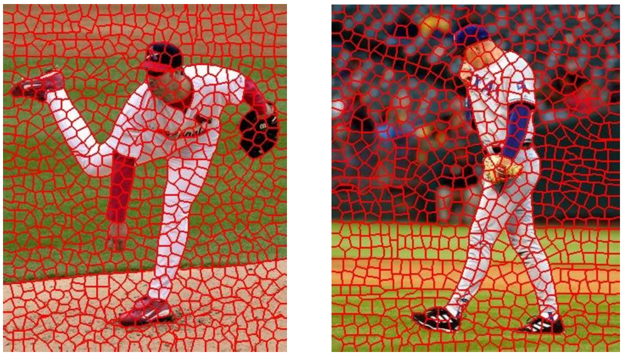

X. Ren and J. Malik. Learning a classification model for segmentation. ICCV 2003.

Source： S. Lazebnik

The goals of segmentation

- Separate image into coherent “objects”
  - “Bottom-up” or “top-down” process?
  - Supervised or unsupervised?

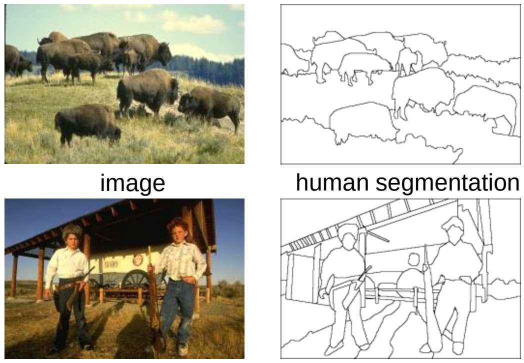

Berkeley segmentation database:
http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/

Inspiration from psychology

- The Gestalt school: Grouping is key to visual perception

The Muller-Lyer illusion

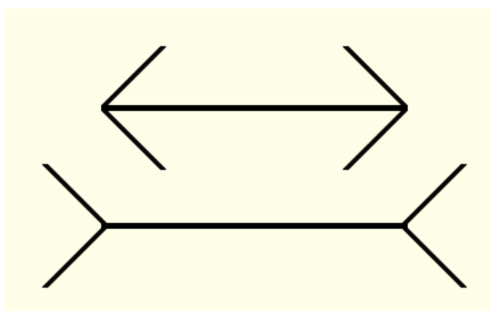

http://en.wikipedia.org/wiki/Gestalt_psychology

The Gestalt school

- Elements in a collection can have properties that result from relationships
  - “The whole is greater than the sum of its parts”

subjective contours

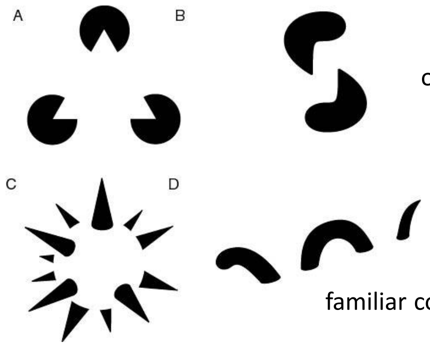

occlusion

familiar configuration

http://en.wikipedia.org/wiki/Gestalt_psychology

Source： S. Lazebnik

Emergence

http://en.wikipedia.org/wiki/Gestalt_psychology

Gestalt factors

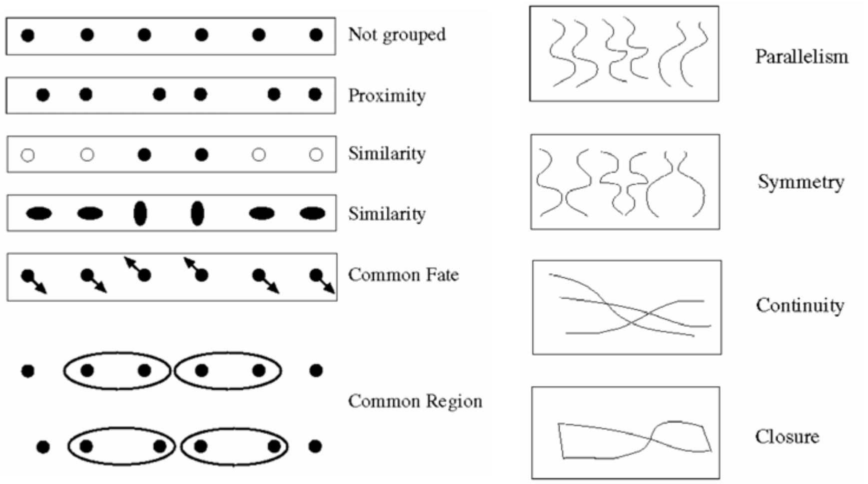

Source： S. Lazebnik

Grouping phenomena in real life

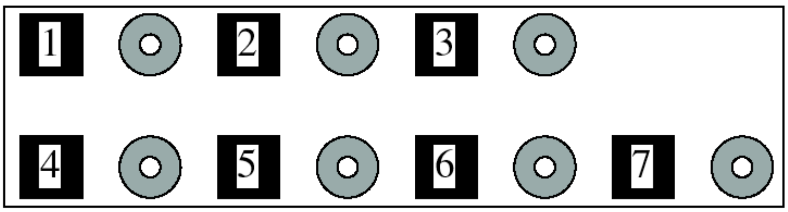

Forsyth & Ponce, Figure 14.7

Source： S. Lazebnik

Grouping phenomena in real life

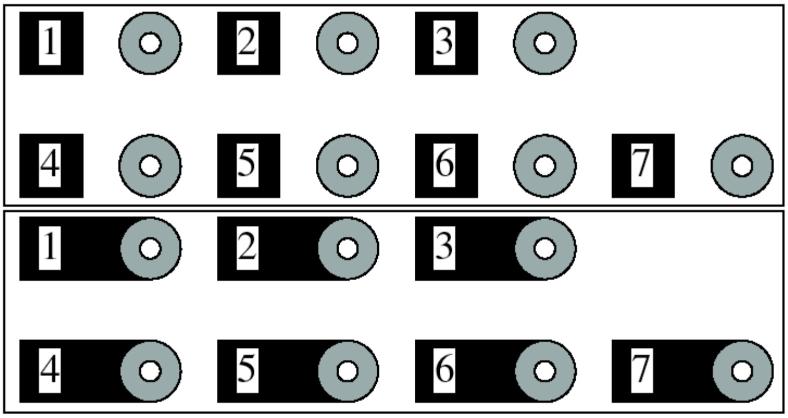

Forsyth & Ponce, Figure 14.7

Source： S. Lazebnik

Gestalt factors

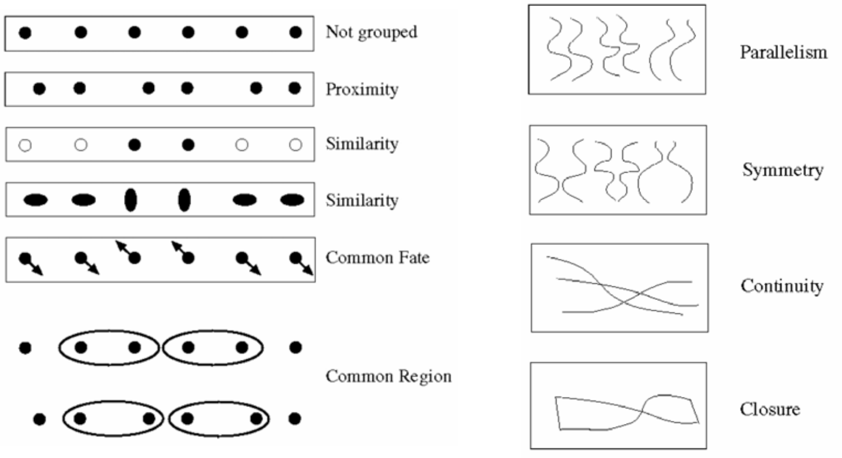

These factors make intuitive sense, but are very difficult to translate into algorithms

Source： S. Lazebnik

Segmentation as clustering

- Cluster similar pixels (features) together

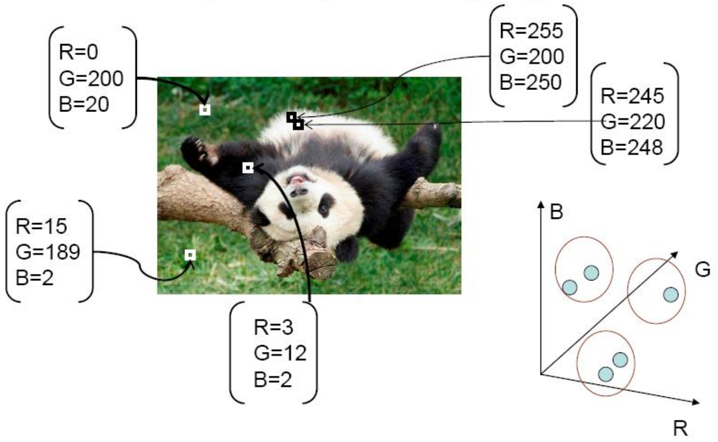

Source: K. Grauman

Segmentation as clustering

- K-means clustering based on intensity or color is essentially vector quantization of the image attributes
  - Clusters don't have to be spatially coherent

Image Intensity-based clusters Color-based clusters

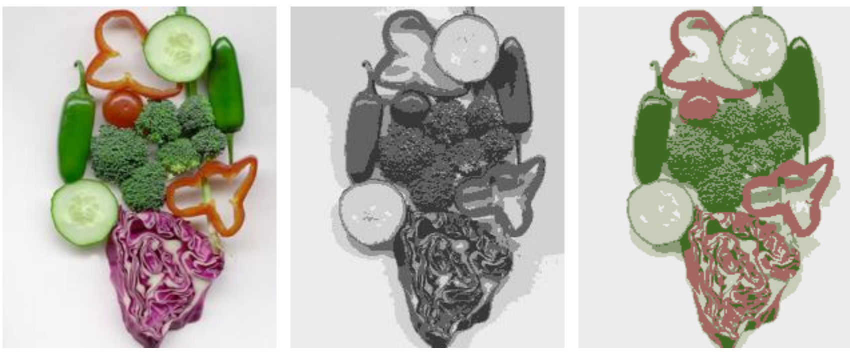

Source： S. Lazebnik

Segmentation as clustering

- Cluster similar pixels (features) together

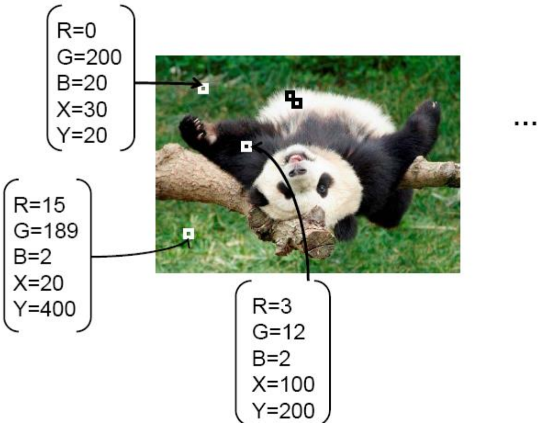

Source: K. Grauman

Segmentation as clustering

- Clustering based on (r,g,b,x,y) values enforces more spatial coherence

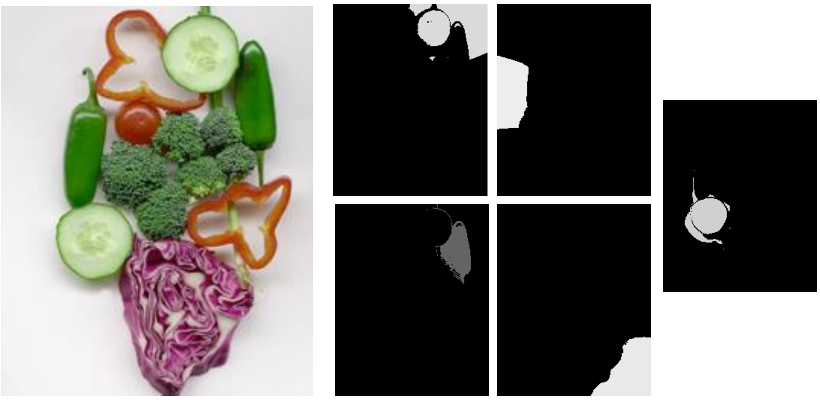

Source： S. Lazebnik

K-Means for segmentation

- Pros
  - Very simple method
  - Converges to a local minimum of the error function
- Cons
  - Memory-intensive
  - Need to pick K
  - Sensitive to initialization
  - Sensitive to outliers
  - Only finds "spherical" clusters

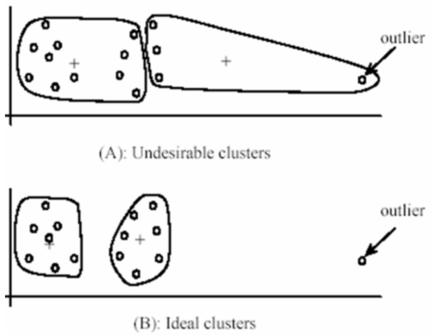

Source： S. Lazebnik

Mean shift clustering and segmentation

- An advanced and versatile technique for clustering-based segmentation

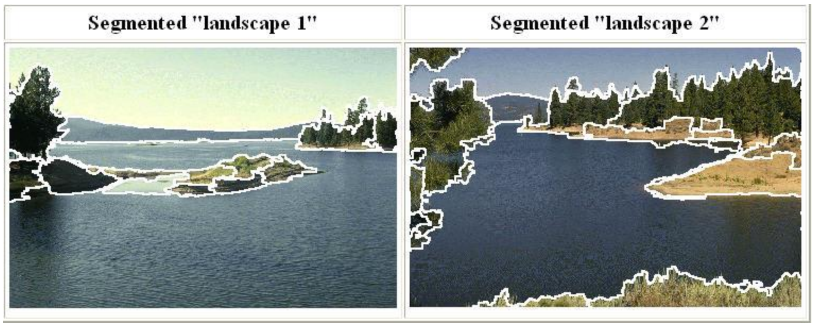

http://www.caip.rutgers.edu/~comanici/MSPAMI/msPamiResults.html

D. Comaniciu and P. Meer, Mean Shift: A Robust Approach toward Feature Space Analysis, PAMI 2002.

Mean shift algorithm

- The mean shift algorithm seeks *modes* or local maxima of density in the feature space

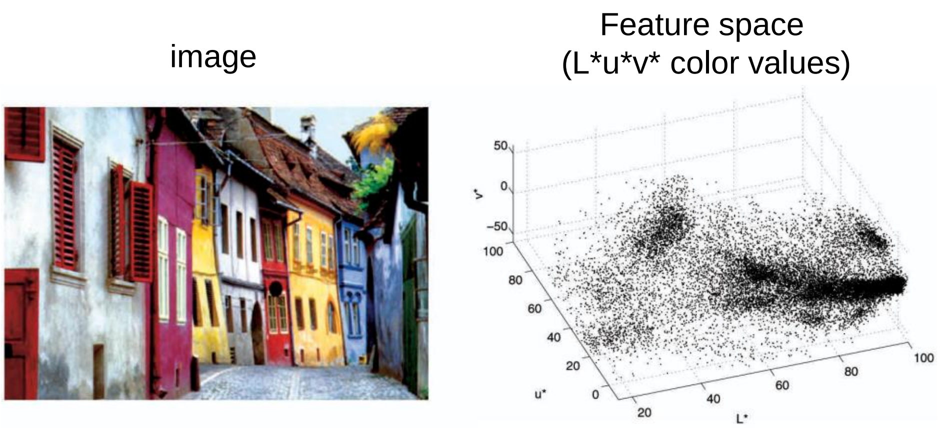

Source：S. Lazebnik

Mean shift

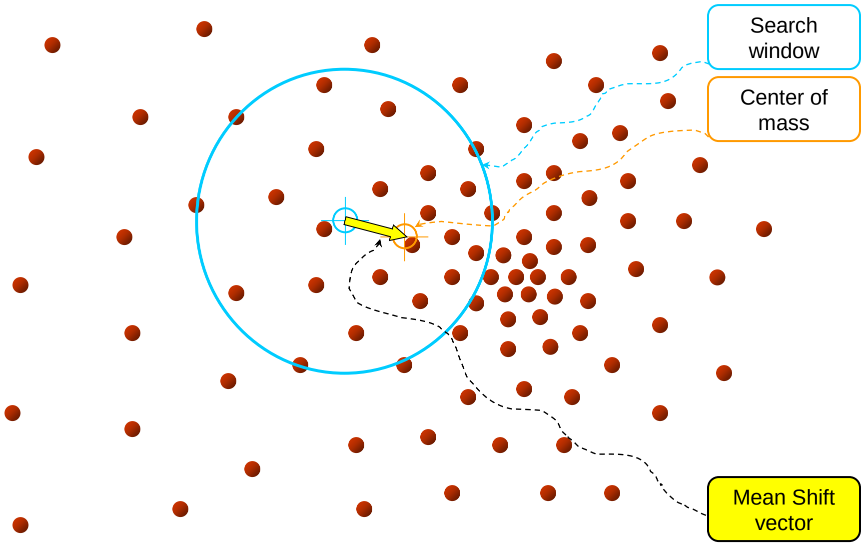

Soruce: Y. Ukrainitz & B. Sarel

Mean shift

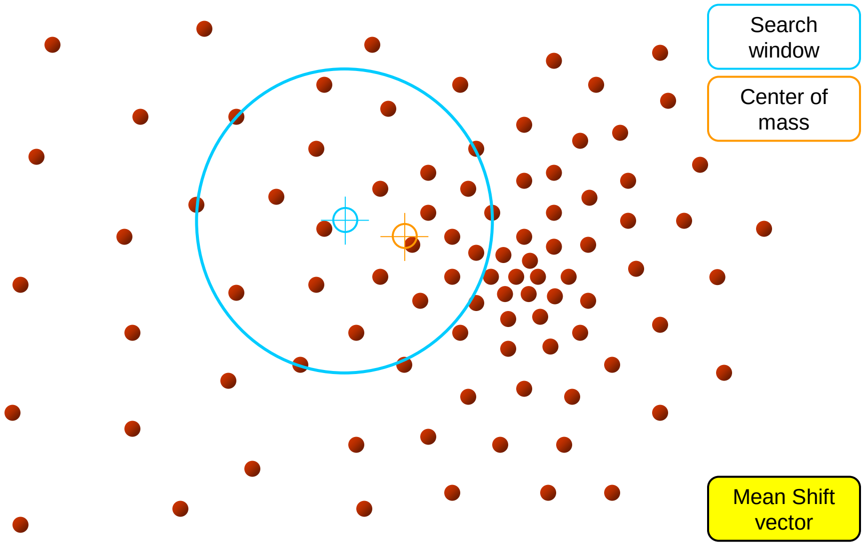

Soruce: Y. Ukrainitz & B. Sarel

Mean shift

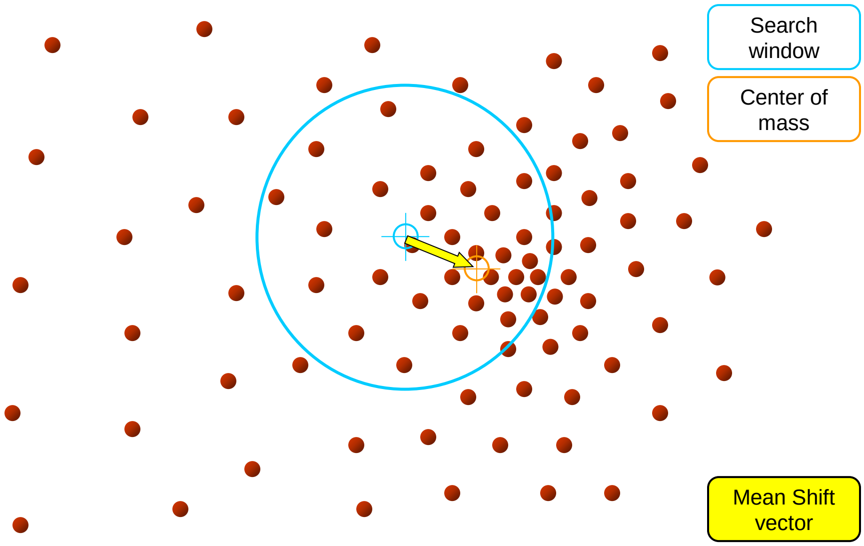

Soruce: Y. Ukrainitz & B. Sarel

Mean shift

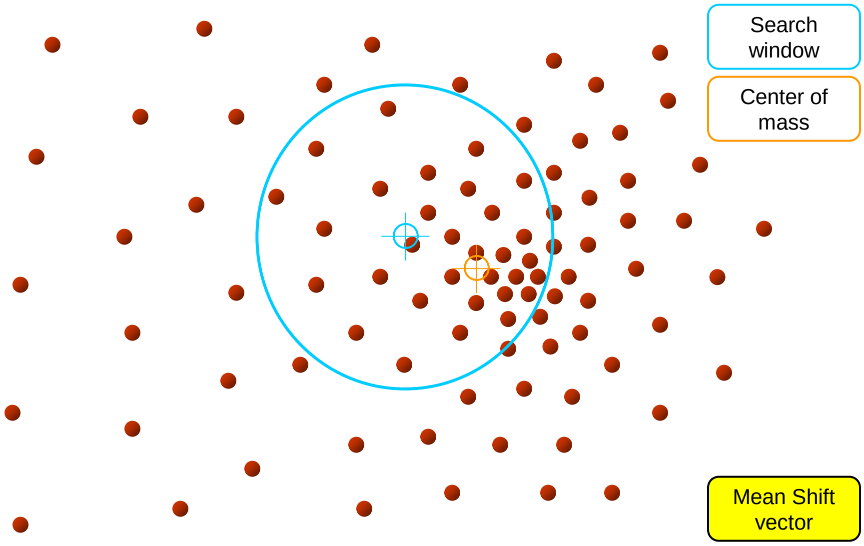

Soruce: Y. Ukrainitz & B. Sarel

Mean shift

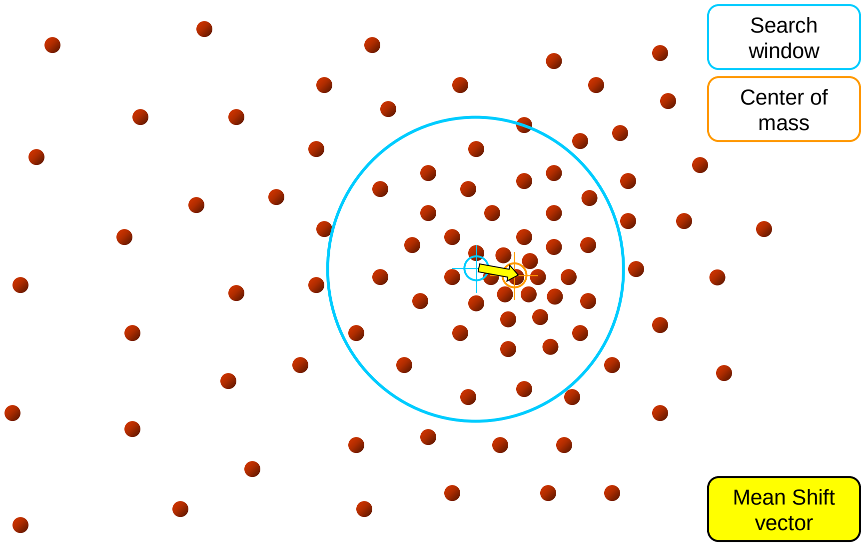

Soruce: Y. Ukrainitz & B. Sarel

Mean shift

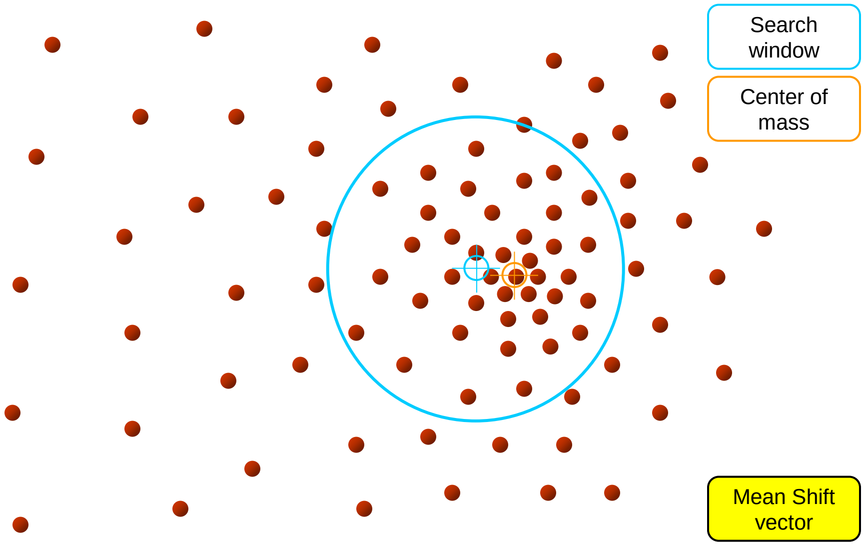

Soruce: Y. Ukrainitz & B. Sarel

Mean shift

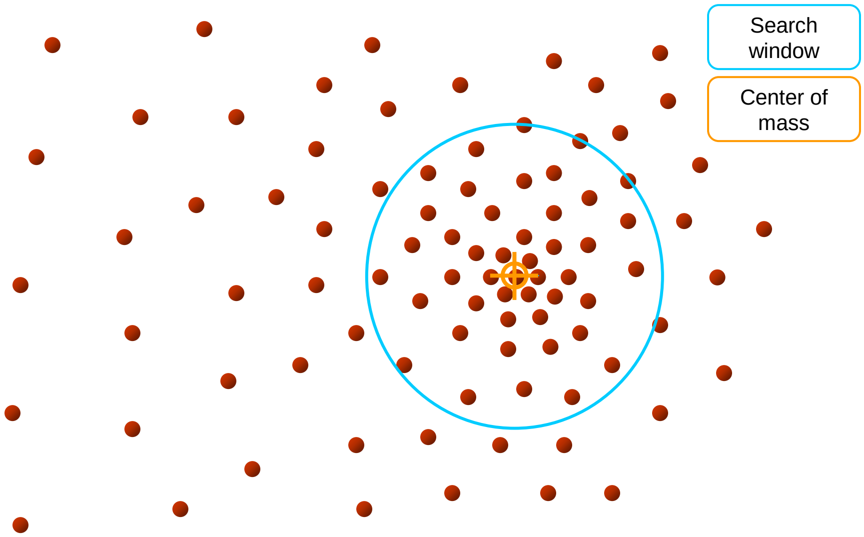

Soruce: Y. Ukrainitz & B. Sarel

Mean shift clustering

- Cluster: all data points in the attraction basin of a mode
- Attraction basin: the region for which all trajectories lead to the same mode

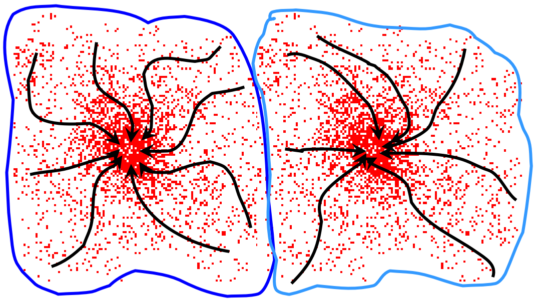

Soruce: Y. Ukrainitz & B. Sarel

Mean shift clustering/segmentation

- Find features (color, gradients, texture, etc)
- Initialize windows at individual feature points
- Perform mean shift for each window until convergence
- Merge windows that end up near the same "peak" or mode

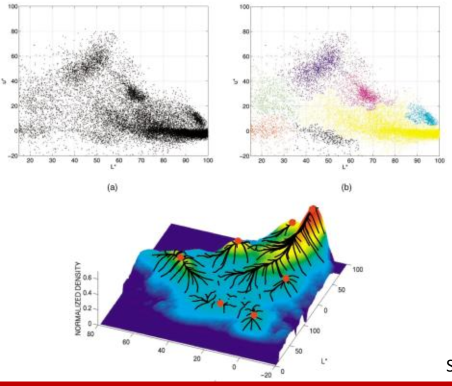

Source： S. Lazebnik

Mean shift segmentation results

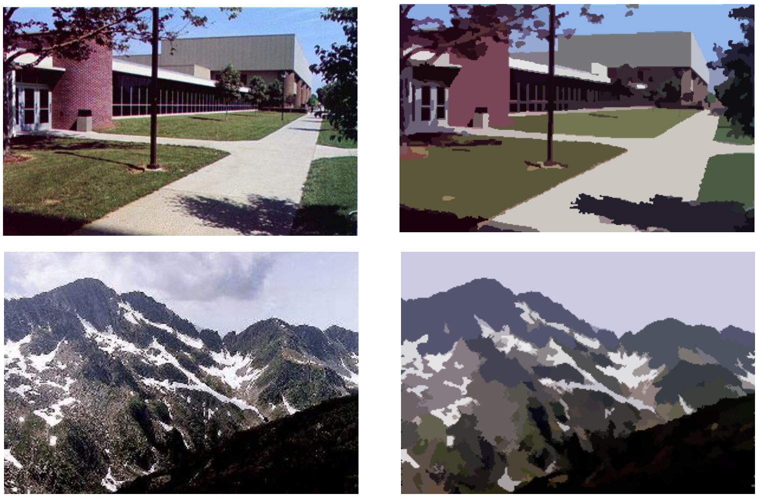

http://www.caip.rutgers.edu/~comanici/MSPAMI/msPamiResults.html

Source： S. Lazebnik

More results

Source： S. Lazebnik

More results

Source： S. Lazebnik

Mean shift pros and cons

- Pros
  - Does not assume spherical clusters
  - Just a single parameter (window size)
  - Finds variable number of modes
  - Robust to outliers
- Cons
  - Output depends on window size
  - Computationally expensive
  - Does not scale well with dimension of feature space

Source： S. Lazebnik

Images as graphs

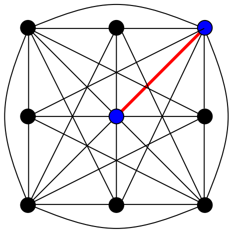

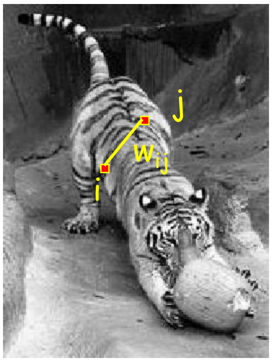

- Node for every pixel
- Edge between every pair of pixels (or every pair of "sufficiently close" pixels)
- Each edge is weighted by the *affinity* or similarity of the two nodes

Source: S. Seitz

Segmentation by graph partitioning

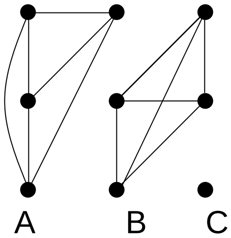

- Break Graph into Segments
  - Delete links that cross between segments
  - Easiest to break links that have low affinity
    - similar pixels should be in the same segments
    - dissimilar pixels should be in different segments

Source: S. Seitz

Measuring affinity

- Suppose we represent each pixel by a feature vector **x**, and define a distance function appropriate for this feature representation
- Then we can convert the distance between two feature vectors into an affinity with the help of a generalized Gaussian kernel:

$$\exp\left(-\frac{1}{2\sigma^2}\operatorname{dist}(\mathbf{x}_i,\mathbf{x}_j)^2\right)$$

Source： S. Lazebnik

Scale affects affinity

- Small σ: group only nearby points
- Large σ: group far-away points

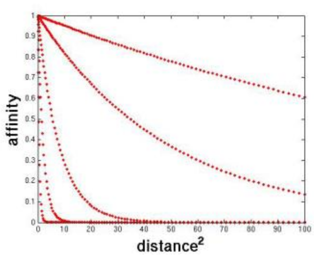

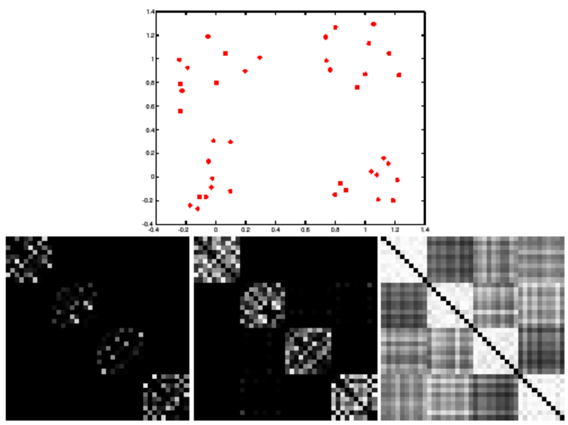

Source： S. Lazebnik

Graph cut

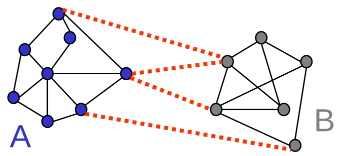

- Set of edges whose removal makes a graph disconnected
- Cost of a cut: sum of weights of cut edges
- A graph cut gives us a segmentation
  - What is a "good" graph cut and how do we find one?

Source: S. Seitz

Minimum cut

- We can do segmentation by finding the *minimum cut* in a graph
  - Efficient algorithms exist for doing this

Minimum cut example

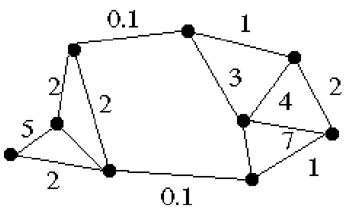

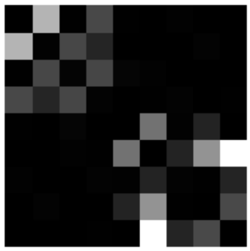

Source： S. Lazebnik

Minimum cut

- We can do segmentation by finding the *minimum cut* in a graph
  - Efficient algorithms exist for doing this

Minimum cut example

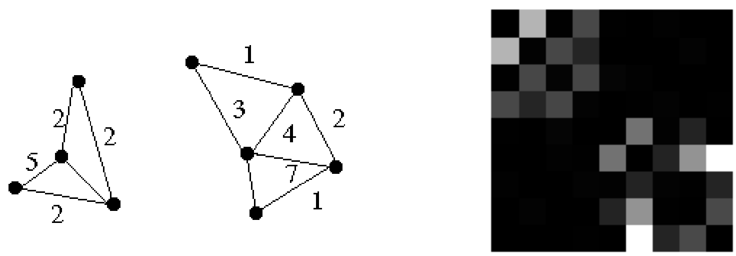

Source： S. Lazebnik

Normalized cut

- Drawback: minimum cut tends to cut off very small, isolated components

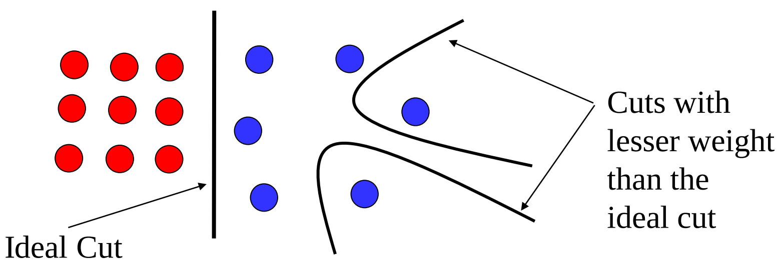

Source： Khurram Hassan-Shafique

Normalized cut

- Drawback: minimum cut tends to cut off very small, isolated components
- This can be fixed by normalizing the cut by the weight of all the edges incident to the segment
- The *normalized cut* cost is:

$$\frac{w(A,B)}{w(A,V)} + \frac{w(A,B)}{w(B,V)}$$

$w(A, B)$ = sum of weights of all edges between A and B

J. Shi and J. Malik. Normalized cuts and image segmentation. PAMI 2000

Source： S. Lazebnik

Normalized cut

- Let W be the adjacency matrix of the graph
- Let D be the diagonal matrix with diagonal entries $D(i, i) = \Sigma_j W(i, j)$
- Then the normalized cut cost can be written as

$$\frac{y^T(D-W)y}{y^T Dy}$$

where y is an indicator vector whose value should be 1 in the ith position if the ith feature point belongs to A and a negative constant otherwise

J. Shi and J. Malik. Normalized cuts and image segmentation. PAMI 2000

Source： S. Lazebnik

Normalized cut

- Finding the exact minimum of the normalized cut cost is NP-complete, but if we *relax* y to take on arbitrary values, then we can minimize the relaxed cost by solving the *generalized eigenvalue problem* $(D - W)y = \lambda Dy$
- The solution y is given by the generalized eigenvector corresponding to the second smallest eigenvalue
- Intutitively, the ith entry of y can be viewed as a "soft" indication of the component membership of the ith feature
  - Can use 0 or median value of the entries as the splitting point (threshold), or find threshold that minimizes the Ncut cost

Source： S. Lazebnik

Normalized cut algorithm

1. Represent the image as a weighted graph $G = (V,E)$, compute the weight of each edge, and summarize the information in D and W
2. Solve $(D - W)y = \lambda Dy$ for the eigenvector with the second smallest eigenvalue
3. Use the entries of the eigenvector to bipartition the graph

To find more than two clusters:

- Recursively bipartition the graph
- Run k-means clustering on values of several eigenvectors

Source： S. Lazebnik

Example result

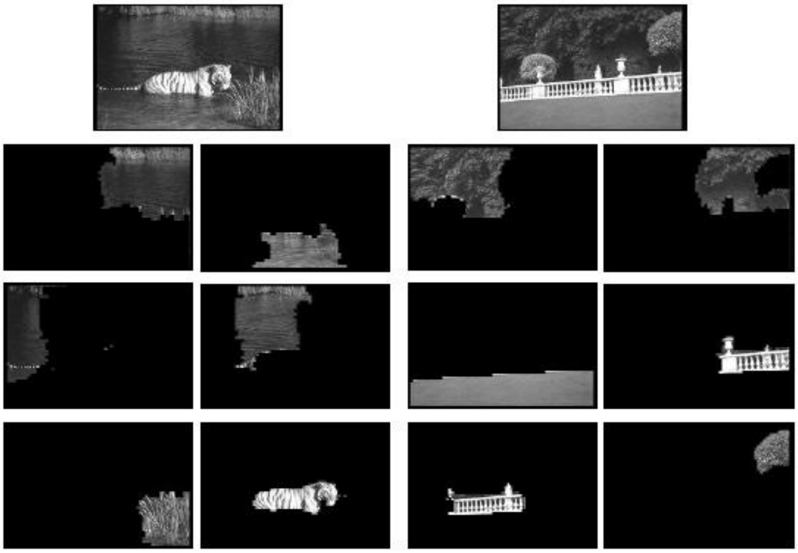

Source： S. Lazebnik

Challenge

- How to segment images that are a “mosaic of textures”?


Source： S. Lazebnik

Using texture features for segmentation

- Convolve image with a bank of filters

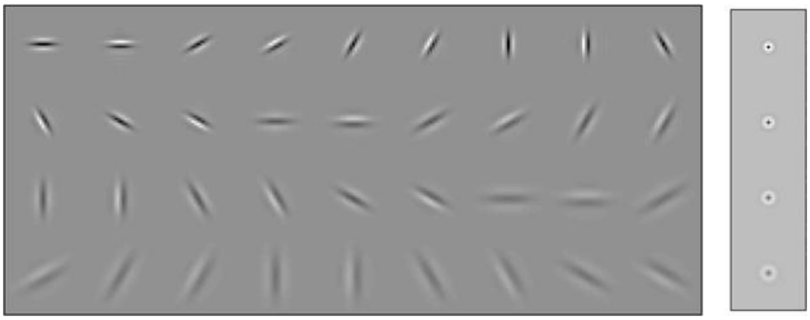

J. Malik, S. Belongie, T. Leung and J. Shi. "*Contour and Texture Analysis for Image Segmentation*". IJCV 43(1),7-27,2001.

Using texture features for segmentation

- Convolve image with a bank of filters
- Find *textons* by clustering vectors of filter bank outputs

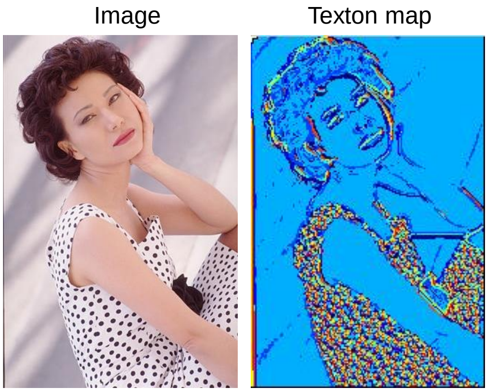

J. Malik, S. Belongie, T. Leung and J. Shi. "*Contour and Texture Analysis for Image Segmentation*". IJCV 43(1),7-27,2001.

Using texture features for segmentation

- Convolve image with a bank of filters
- Find *textons* by clustering vectors of filter bank outputs
- The final texture feature is a texton histogram computed over image windows at some "local scale"

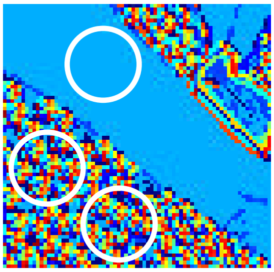

J. Malik, S. Belongie, T. Leung and J. Shi. "*Contour and Texture Analysis for Image Segmentation*". IJCV 43(1),7-27,2001.

Pitfall of texture features

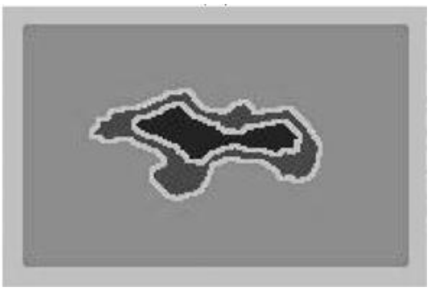

- Possible solution: check for "intervening contours" when computing connection weights

J. Malik, S. Belongie, T. Leung and J. Shi. "*Contour and Texture Analysis for Image Segmentation*". IJCV 43(1),7-27,2001.

Example results

Source： S. Lazebnik

Results: Berkeley Segmentation Engine

http://www.cs.berkeley.edu/~fowlkes/BSE/

Source： S. Lazebnik

Normalized cuts: Pro and con

- Pros
  - Generic framework, can be used with many different features and affinity formulations
- Cons
  - High storage requirement and time complexity
  - Bias towards partitioning into equal segments

Source： S. Lazebnik

Segments as primitives for recognition

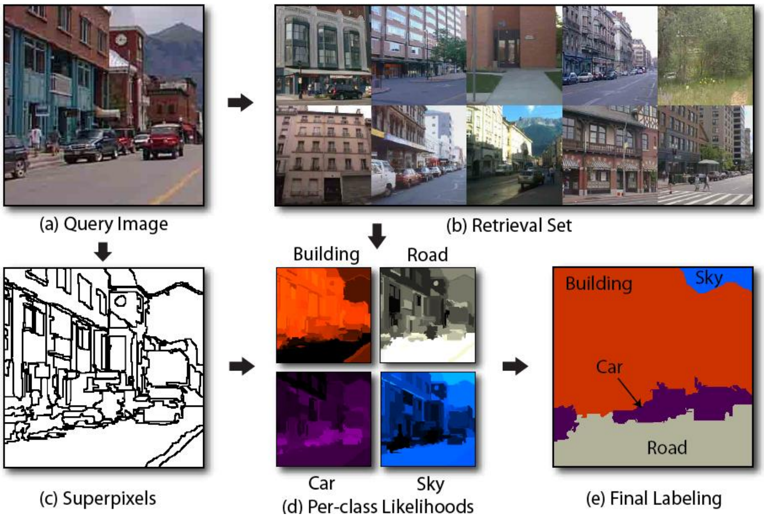

J. Tighe and S. Lazebnik, ECCV 2010